

Part 5: Working in RStudio.Cloud

In this chapter, we're going to learn about R packages, navigating the RStudio panes, some additional benefits of working in Rmarkdown and plain text, and how to make and monitor changes with Github.

Navigating the panes in our workbench

The next few sections will walk through a few of RStudio.Cloud's panes. To recap:

1) *We uploaded files into the `project` folder, which now has the following structure:*

```
project/
├── CHANGELOG.txt
├── README.Rmd
├── README.md
├── code/
├── data/
├── dem-pres-debate-2019.Rproj
├── figs/
└── project.Rproj
```

2) *There are three folders in this project: `code/`, `data/`, and `figs/`.*

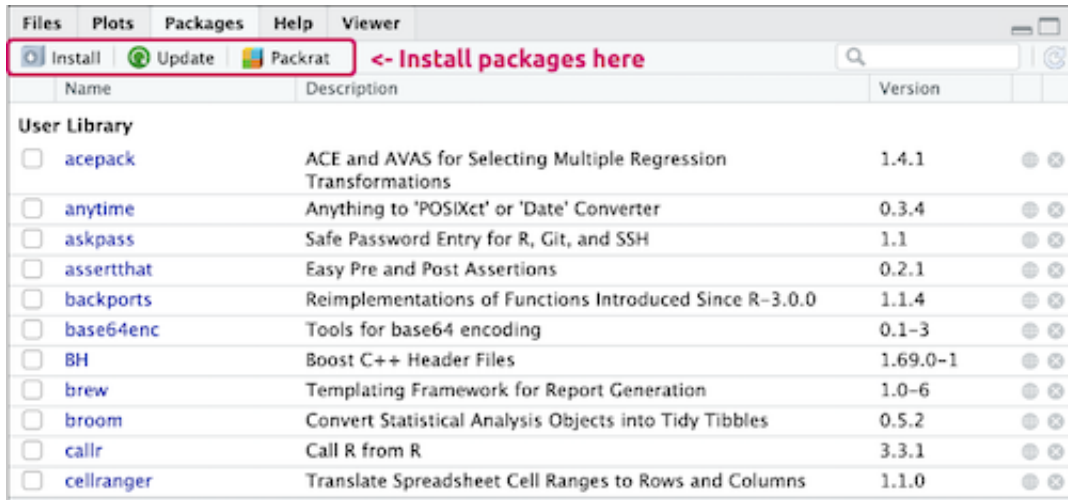
3) *There are three types of plain text files in the parent folder: `.md`, `.Rmd` and `.txt`, and two project files: `dem-pres-debate-2019.Rproj` and `project.Rproj`.*

We'll go over these folders and files in more depth in the following sections.

Packages

Packages are a vital part of the R ecosystem. R packages are collections of functions and objects collected together with a specific purpose. When we started our RStudio.Cloud session, a few packages get loaded automatically. R users typically refer to these default packages and functions as “base R.” Base R packages cover a wide range of statistical procedures and visualizations. Unfortunately, base R can also be challenging to learn because of its inconsistent syntax and style conventions.

A list of the packages and their descriptions are available in the **Packages** pane. We've installed the packages we will use with the `00.1-inst-packages.R` file. If we wanted to, we could also install more packages using *Install* icon, and then enter its name. We don't recommend this, though, because you'll want a record of the packages you used in the project.



You can click on the names of each package to learn more about them and load them into the RStudio session. One of the great things about R is the many user-created packages that greatly expand the number of functions. At the time of this writing, R users have contributed [14638](#) packages available for us to download and use.

The tidyverse

R is open-source software, so users can write packages to expand and enhance its functionality. The [tidyverse](#) is a collection of R packages from RStudio for doing data science. All `tidyverse` packages share a few similar underlying principles that allow them to work well together.

Unlike base R, the `tidyverse` also has a consistent grammar and syntax, which makes it easier to read and write. You can learn more about this syntax in the [R for Data Science](#) text or on the [tidyverse webpage](#).

So far, the code we've run comes from base R. Going forward; we're going to use various packages from the `tidyverse`.

Files

The **Files** pane displays the files and folders in this project. The file path is visible in the top portion of the pane, beneath the options for *New Folder*, *Upload*, *Delete*, *Rename*, and *More*. An image of this folder's contents is below:

Files	Plots	Packages	Help	Viewer
New Folder Upload Delete Rename More ▾				
<input type="checkbox"/> Cloud > project > code				
▲ Name				Size
..				
<input type="checkbox"/>		00.1-inst-packages.R		768 B
<input type="checkbox"/>		00.2-download-538.R		2.7 KB
<input type="checkbox"/>		00.3-download-google.R		2.5 KB
<input type="checkbox"/>		00.4-download-tweets.R		3.4 KB
<input type="checkbox"/>		00.5-download-wikipedia.R		2.4 KB
<input type="checkbox"/>		01-import.R		2.9 KB
<input type="checkbox"/>		02-wrangle.R		9.4 KB

The script to install the packages for this project (00.1-inst-packages.R) is in the code folder. Navigate to this folder either by clicking on the file path (Cloud/project/code).

We can see 00.1-inst-packages.R and other scripts are in the **Files** pane. We now know that because the .R files are plain text files, so they have code for the computer to execute, and comments for a human to read.

In R, we can create comments with a preceding # on any line.

```
# this is a comment
```

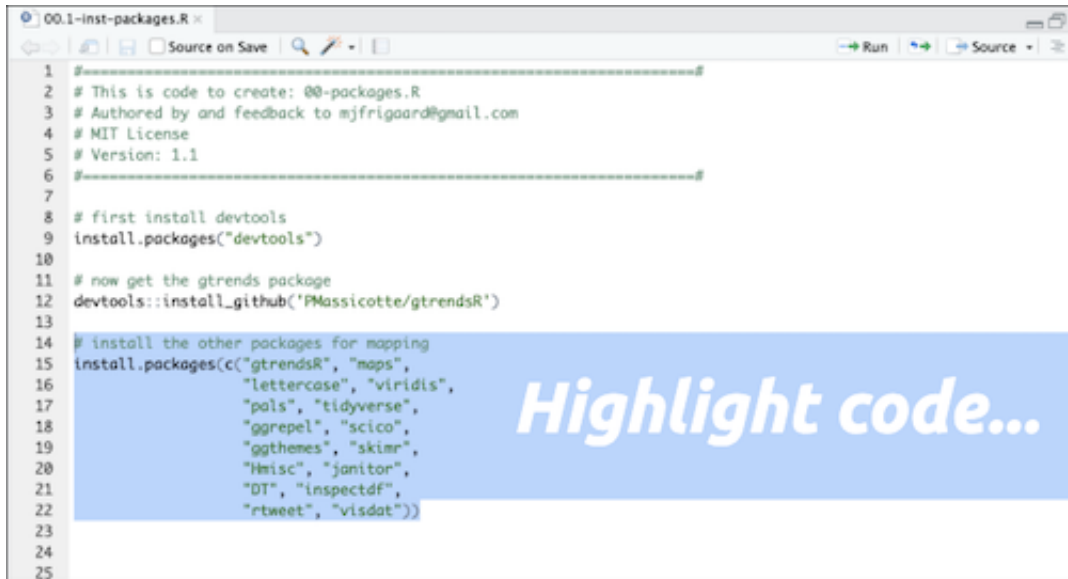
When we click on the 00.1-inst-packages.R file, we can see it open in the **Source** pane. RStudio.Cloud is giving us a few hints about the packages referenced in the .R file. We're being told we need to install a package before moving forward:

Package devtools required but is not installed.

RStudio gives us a choice to ***Install*** or ***Don't Show Again***, and we'll click on the ***Install*** option. Packages vary in the length they take to install, but we'll wait patiently for devtools package finish downloading. After the install has finished, you should see the > prompt in the **Console** pane.

Source

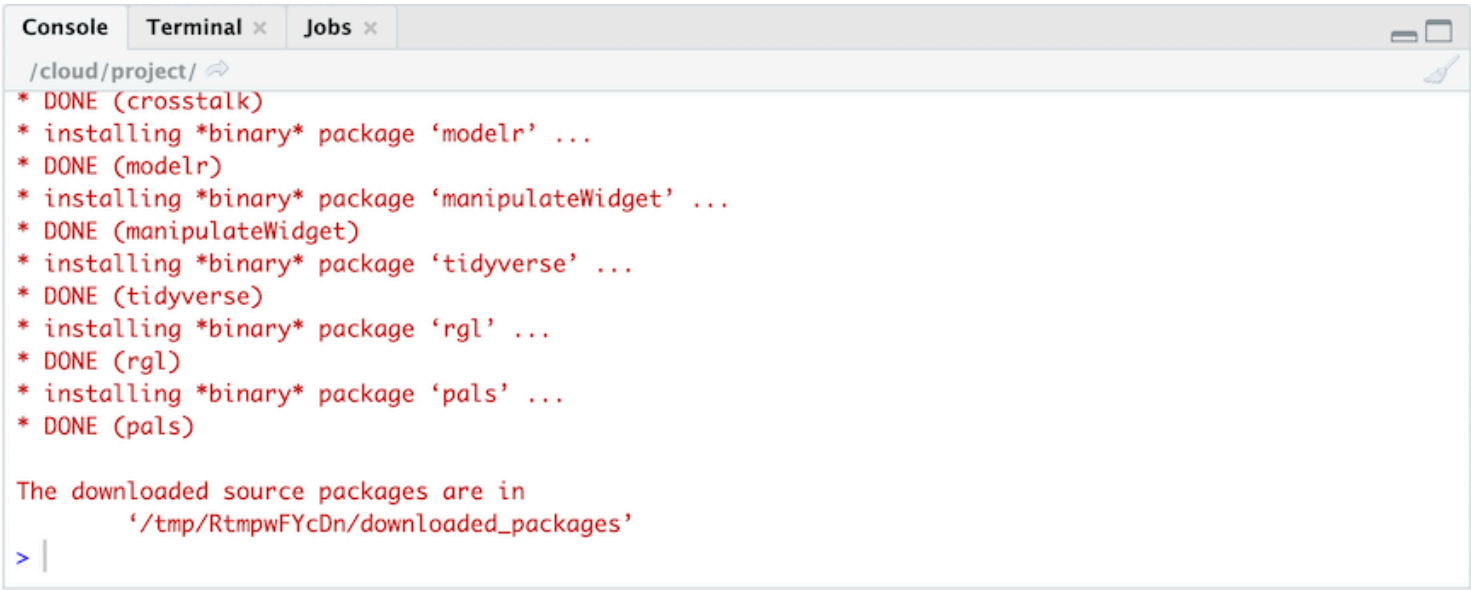
When working in RStudio, sometimes we'll write some code, and want to run that code and see if it works. To do this, we can highlight the rest of the script with the mouse cursor (lines 14–22), then hold down **ctrl** or **cmd** and hit **enter** or **return**.



```
1 #-----
2 # This is code to create: 00-packages.R
3 # Authored by and feedback to mjfrigaard@gmail.com
4 # MIT License
5 # Version: 1.1
6 #-----
7
8 # first install devtools
9 install.packages("devtools")
10
11 # now get the gtrends package
12 devtools::install_github("PMassicotte/gtrendsR")
13
14 # install the other packages for mapping
15 install.packages(c("gtrendsR", "maps",
16                   "lettercase", "viridis",
17                   "pals", "tidyverse",
18                   "ggrepel", "scico",
19                   "ggthemes", "skmr",
20                   "Hmisc", "janitor",
21                   "DT", "inspectdf",
22                   "rtweet", "visdat"))
23
24
25
```



After the installation has completed for all the packages, we'll see the following in the **Console** pane.



```
Console Terminal x Jobs x
/cloud/project/
* DONE (crosstalk)
* installing *binary* package 'modelr' ...
* DONE (modelr)
* installing *binary* package 'manipulateWidget' ...
* DONE (manipulateWidget)
* installing *binary* package 'tidyverse' ...
* DONE (tidyverse)
* installing *binary* package 'rgl' ...
* DONE (rgl)
* installing *binary* package 'pals' ...
* DONE (pals)

The downloaded source packages are in
  '/tmp/RtmpwFYcDn/downloaded_packages'
> |
```

Back in the **Files** pane, we can navigate back to the `project/` folder by clicking on the path above the files.

Files Plots Packages Help Viewer

New Folder Upload Delete Rename More

Cloud > project > code

Click here

Name	Size	Modified
..		
00.1-inst-packages.R	768 B	Jul 28, 2019, 2:07 PM
00.2-download-538.R	2.7 KB	Jul 26, 2019, 9:20 AM
00.3-download-google.R	2.5 KB	Jul 26, 2019, 9:20 AM
00.4-download-tweets.R	3.4 KB	Jul 26, 2019, 9:20 AM
00.5-download-wikipedia.R	2.4 KB	Jul 26, 2019, 9:20 AM
01-import.R	2.9 KB	Jul 26, 2019, 9:20 AM
02-wrangle.R	9.4 KB	Jul 26, 2019, 9:20 AM

We can see the `README.Rmd` file in the **Files** pane in the lower right corner. We'll click on it, so the file opens in the **Source** pane (see image below).

File Edit Code View Plots Session Build Debug Profile Tools Help

R 3.6.0

Environment History Connections Git

Diff Commit Pull Push History More

Staged Status Path

Staged	Status	Path
		.gitignore
		01-import.Rmd
		02-wrangle.Rmd
		03-visualize.Rmd
		CHANGELOG.txt
		README.Rmd
		README.md
		code/
		data/
		dem-pres-debate-2019.Rproj
		figs/
		project.Rproj

Files Plots Packages Help Viewer

New Folder Upload Delete Rename More

Cloud > project

Name	Size	Modified
..	0 B	Jul 26, 2019, 9:20 AM
01-import.Rmd	13.4 KB	Jul 26, 2019, 9:20 AM
02-wrangle.Rmd	10 KB	Jul 26, 2019, 9:20 AM
03-visualize.Rmd	1.3 KB	Jul 26, 2019, 9:20 AM
CHANGELOG.txt	86 B	Jul 26, 2019, 9:33 AM
code		
data		
dem-pres-debate-2019.Rproj	205 B	Jul 26, 2019, 9:20 AM
figs		
project.Rproj	205 B	Jul 26, 2019, 9:38 AM
README.md	16.5 KB	Jul 26, 2019, 9:20 AM
README.Rmd	12.8 KB	Jul 26, 2019, 9:20 AM

Source

```

1 ---
2 title: "2019 Democratic Debates data project"
3 author: "Jane Doe"
4 output: github_document
5 ---
6
7 ```{r setup, include=FALSE}
8 # create image folder ----
9 if (!file.exists("figs/")) {
10   dir.create("figs/")
11 }
12 # create data folder ----
13 if (!file.exists("data/")) {
14   dir.create("data/")
15 }
16 knitr::opts_chunk$set(
17   echo = TRUE, # show all code
18 )
19
20 2019 Democratic Debates data project
  
```

Console

```

* DONE (shades)
* installing *binary* package 'DT' ...
* DONE (DT)
* installing *binary* package 'rtweet' ...
* DONE (rtweet)
* installing *binary* package 'visdat' ...
* DONE (visdat)
* installing *binary* package 'ggfittext' ...
* DONE (ggfittext)
* installing *binary* package 'inspectdf' ...
* DONE (inspectdf)

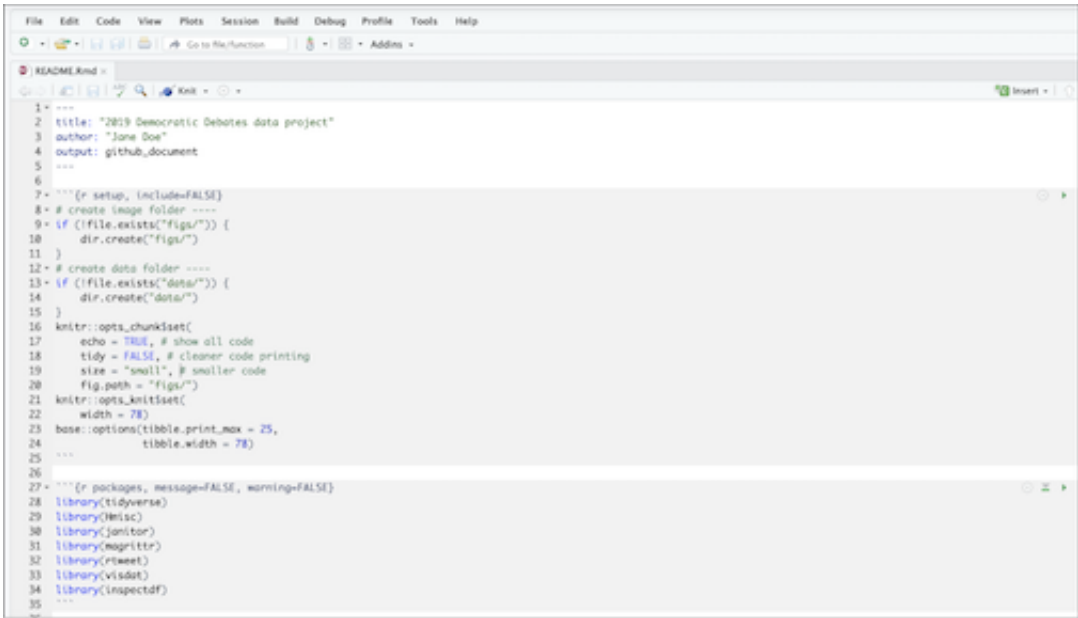
The downloaded source packages are in
  '/tmp/RtmpwFYcDn/downloaded_packages'
  
```

We've highlighted the pane of RStudio that's displaying the `README.Rmd` file. As you can see, this is a relatively small area in the IDE. Working in a tiny corner of the screen can be hard, but fortunately, RStudio gives us the ability to expand any pane into a fullscreen view. In this case, we'll zoom in the `README.Rmd` file.

If we want to focus on the **Source** pane, and can zoom in using `shift+control+1`.



As soon as you click 1, your screen should expand to look like the image below:



Since the **Source** pane is where we'll write most of our code, it's also the pane we are spending most of our time. We've seen that's where our files open, so being able to focus on the area quickly is helpful.

Console

The **Console** is probably the second most used area in RStudio (it displays most of the output), so we also should know how to zoom in on this pane. To focus on this pane, we'll use `shift+control+2`.

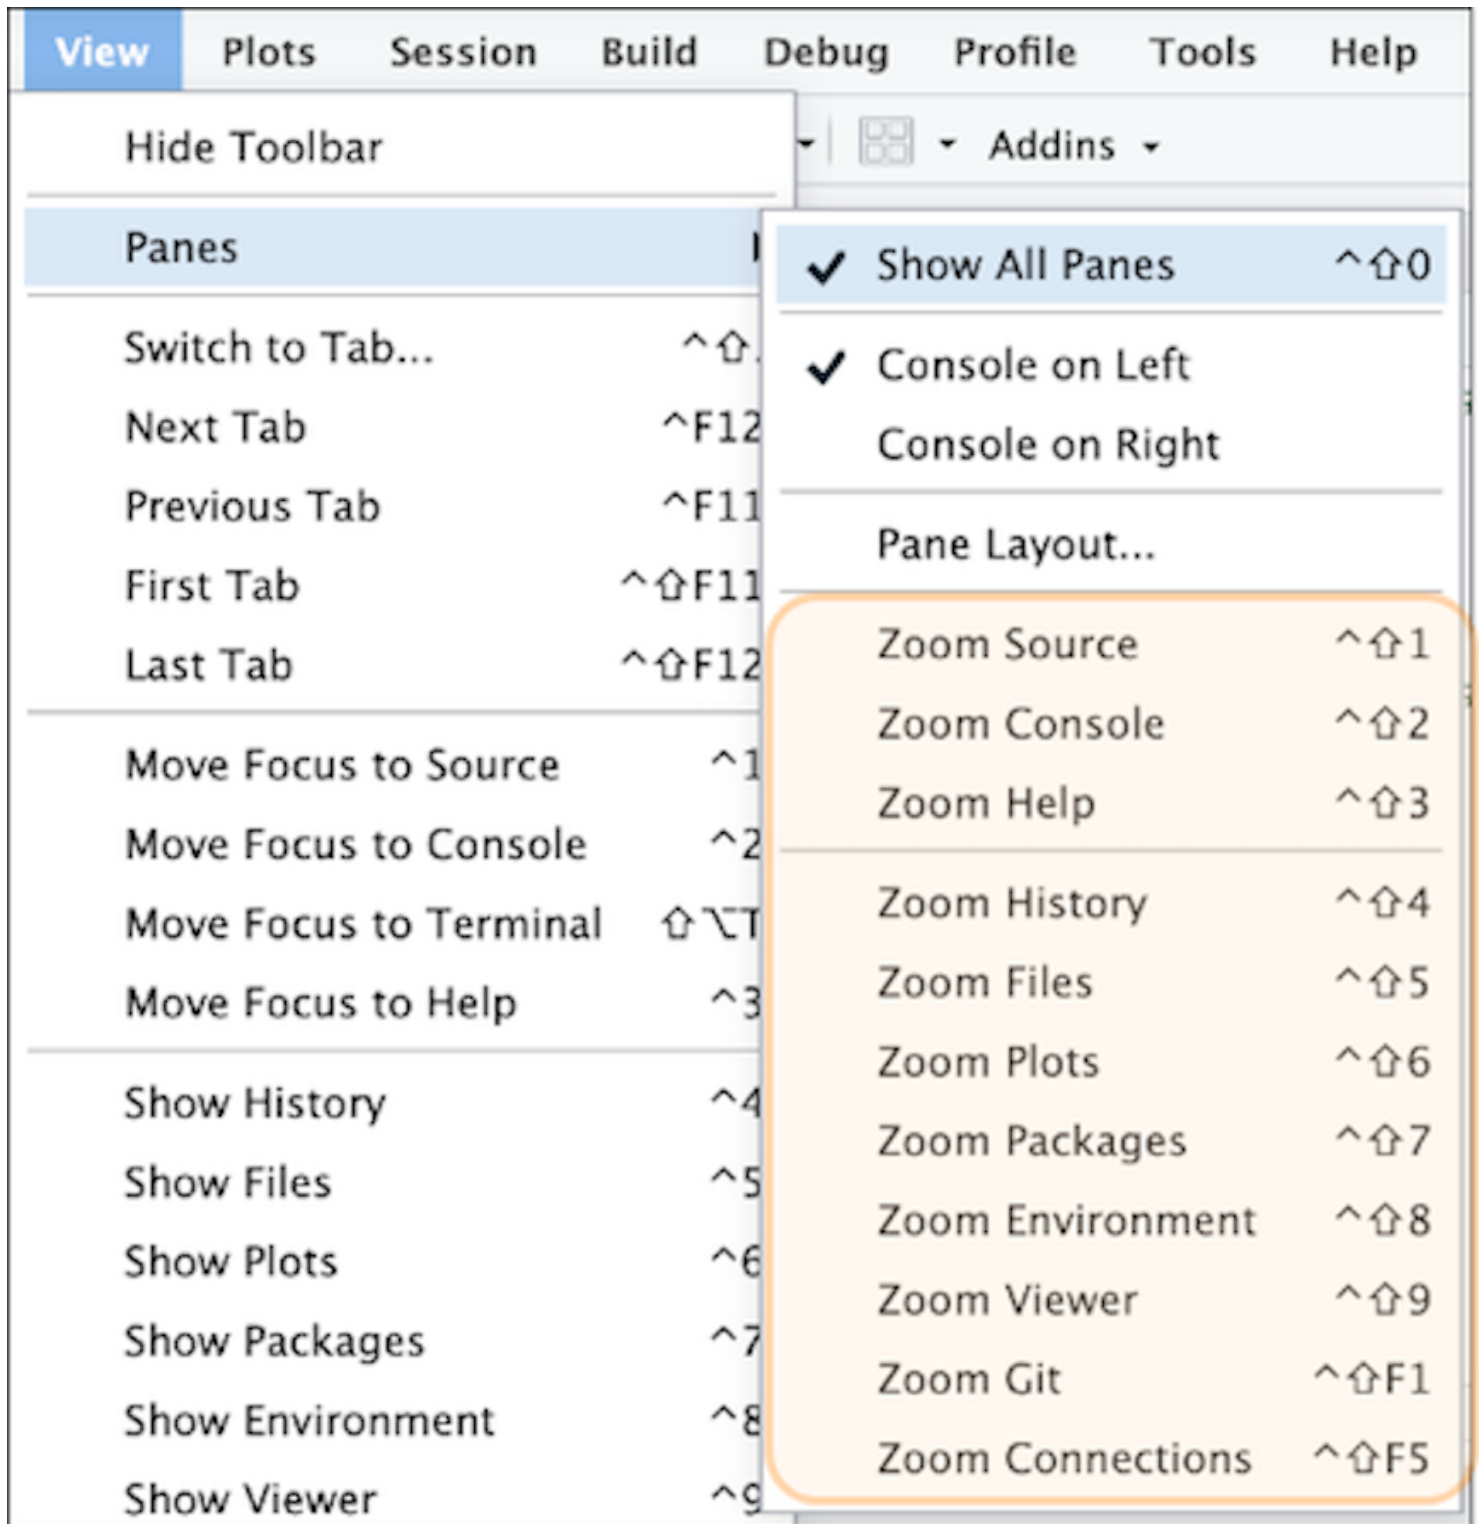


What if we accidentally zoom into the wrong pane? We can easily resize the IDE to its original position by holding down the `shift+control` buttons, then clicking either number again. Combining those keys should return the IDE to its default arrangement.

Working in one pane

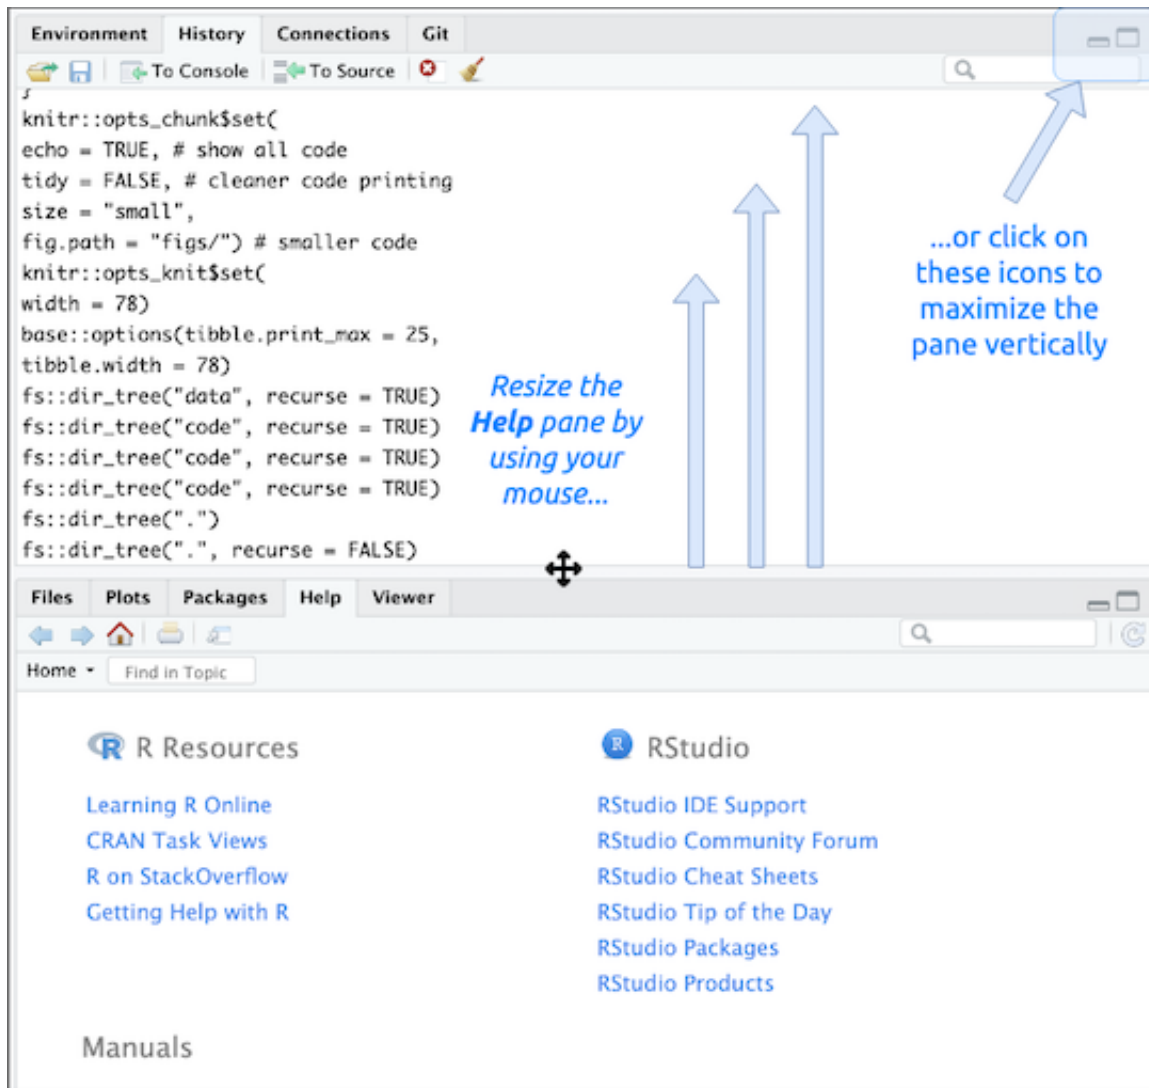
We can move through every pane in RStudio by holding down `shift+control` and clicking on numbers 1-9. Go ahead and do that now.

If you ever forget which number corresponds to which pane, you can always find them under *View > Panes* (see image below)



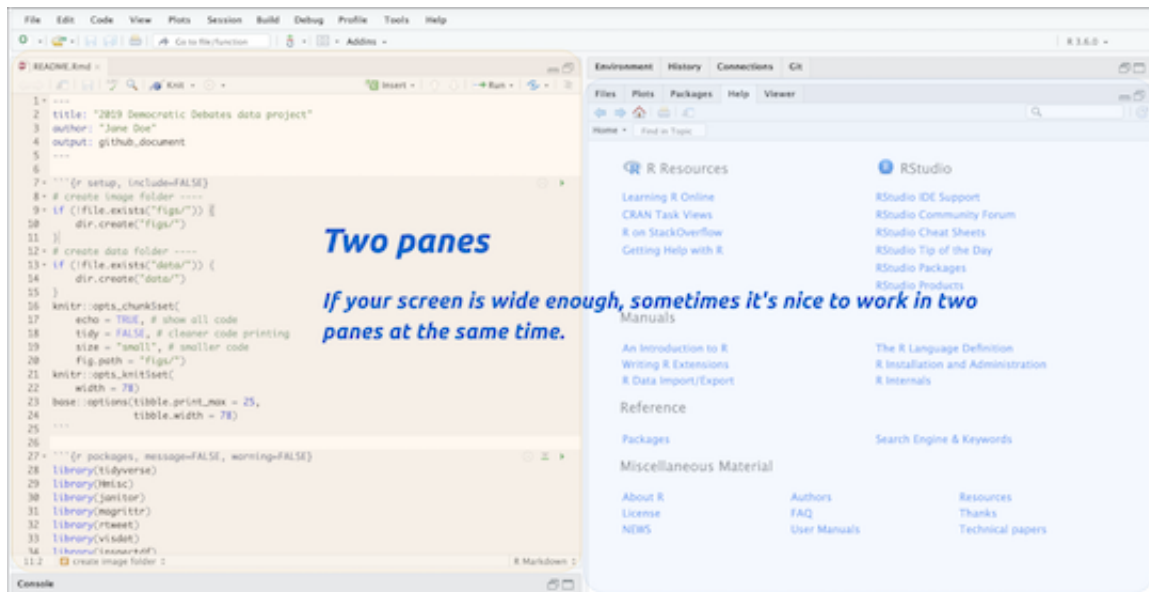
Help

Inevitably, we'll write something that doesn't work. When things aren't working (try to remain calm), you're going to want a place to start looking for solutions. The **Help** pane is accessible in the lower right corner of the IDE.



Working in two panes

Sometimes we'll want to do our work in more than one pane at a time. For example, what if we're working in the **Source** pane, but have a question about a function? We can get answers to a lot of problems using RStudio's internal **Help** pane. After resizing the **Help** pane, we should see the following layout in our IDE.



Recap on RStudio panes

As we pointed out earlier, RStudio is like a workbench built around different panes. Each pane serves a specific purpose and being able to move between them allow us to work quickly and efficiently.

Keyboard shortcuts are time-savers. Not having to switch from keyboard to mouse over and over again keeps us focused on the task at hand.

Where do we write code?

You might be wondering why we have provided you with both `.R` scripts and `.Rmd` files. Well, we want to show you a few options for documenting your work in RStudio. The two sections below outline two standard options (but these are by no means the only way to work with these tools!)

Option 1) everything goes in scripts

`rmarkdown` is a relatively new package, so early R users didn't have an option to put everything in `.Rmd` files (we fit into this group). We have found this approach is excellent when you're reasonably confident about the project. If you're familiar with the data files, wrangling techniques, visualizations, and products, you can quickly outline the scripts and predetermine their names and contents.

For example, the folder tree of our code folder tells us what we should expect from this project.

```
code
├── 00.1-inst-packages.R
├── 00.2-download-538.R
├── 00.3-download-google.R
├── 00.4-download-tweets.R
├── 00.5-download-wikipedia.R
└── 01-import.R
```

We can see there are seven script files, each one with a numerical prefix. These prefixes tell us the order to run the code files, too. The file names also tell us what each script does (download, import, wrangle, etc.).

The layout above is an example of how to organize a set of .R scripts that follow the guidelines mentioned in previous chapters. These naming conventions make it easier for newcomers to the project to get oriented to its contents.

Option 2) everything goes in the Rmarkdown

The second option is well illustrated in the [tweet](#) below by Andrew MacDonald:



Andr(é) MacDonalD
@polesasunder



ok now listen, the harsh truth is you're better off writing one thick, messy .Rmd where you keep all your garbage models and weird musings then going off on some precious folder structure and artfully-named .R files where you can't find a damned thing ever. [#rstats](#) [#oldman](#)

5:45 AM · Jan 17, 2018 · [TweetDeck](#)

<https://twitter.com/polesasunder/status/953624238266646529?s=20>

In our experience, this accurately captures the reality of data projects (especially in the beginning stages). Throwing everything into the Rmarkdown file gives us a lot more flexibility by allowing us to add multiple types of content.

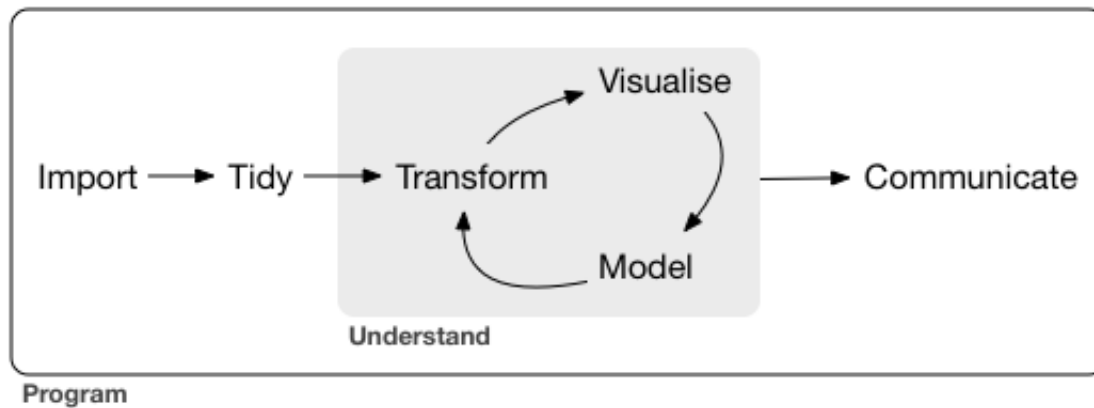
All that being said, it's always good practice to come back and revise the `README.Rmd` document, so it's contents outline all steps in the analysis thoroughly. We consider documentation to be a form of communication with our future selves, and this typically involves clearly describing each step. During the revision, we might also break down each of the chunks into individual scripts. We can also add more details to the text portions, with links to external content, etc.

We've found each project usually starts at a bit of a sprint, so we try to capture as much code and content in .Rmd files early. We can focus on trimming it down later, but the flexibility of markdown combined with functional code is better than limiting ourselves to code and comments .R script files.

Rmarkdown

We're going to move forward assuming we're documenting everything in a `README.Rmd` file. In the next few sections, we'll add various components to the `README.Rmd` we've placed in the project folder.

If you're ever wondering how to outline your `README.Rmd` file, the figure from [R for Data Science](#) isn't a wrong place to start.



.Rmd step 1: the YAML header

At the top of the `README.Rmd` document, the first thing we see is what's called the "YAML header", and it's going to tell RStudio.Cloud the file's title, the author, and what the output file will be.

```
1 ---
2 title: "2019 Democratic Debates data project"
3 author: "Jane Doe"
4 output: github_document
5 ---
```

The YAML header always goes at the top of the `README.Rmd` file, between two sets of three dashes:

```
---
title: "2019 Democratic Debates data project"
author: "Jane Doe"
output: github_document
---
```

People typically use YAML in configuration files, which makes it perfect for setting some default options in our `README.Rmd` document. Read more about YAML headers in the [RMarkdown book](#) and on the [YAML website](#). YAML actually stands for “YAML Ain’t Markup Language.”

We can change the output argument to `html_document`, `word_document`, or `pdf_document` and create a different file from the plain text we are going to be working in. For now, we are going to focus on the `github_document` output.

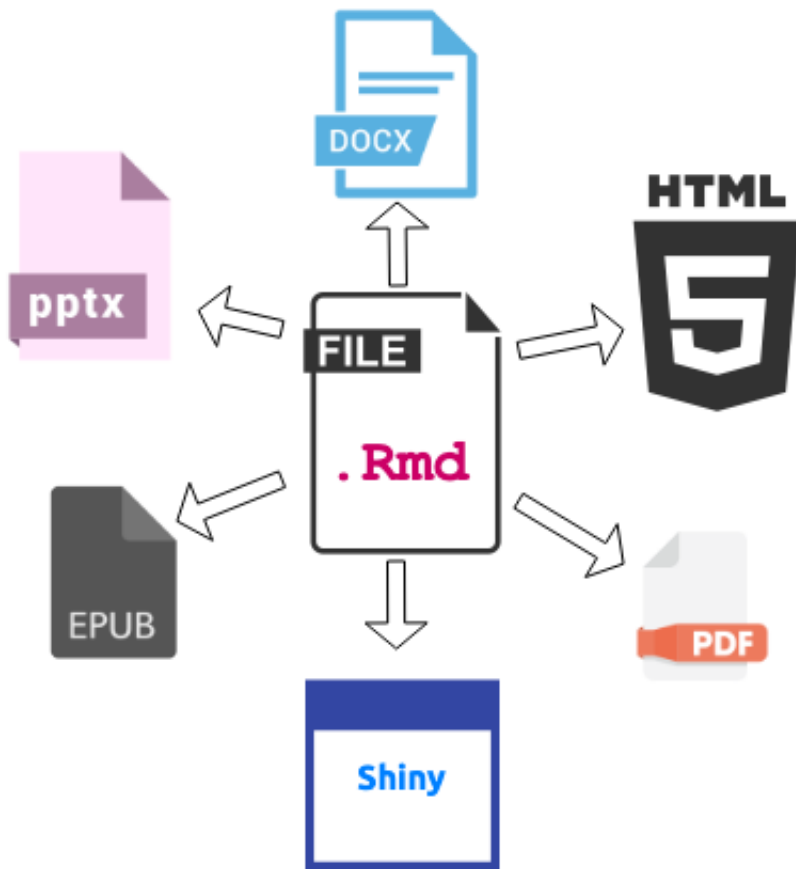
.Rmd step 2: Knit output options

Another benefit of working in the `README.Rmd` document is that when we combine it with the powerful [knitr package](#), we drastically extend the kind of files we can produce from our analysis. `Knitr` follows a principle of [literate programming put forth by Donald E. Knuth](#).

“Let us change our traditional attitude to the construction of programs: Instead of imagining that our main task is to instruct a computer what to do, let us concentrate rather on explaining to humans what we want the computer to do.”

- Donald E. Knuth, *Literate Programming*, 1984

The **Knit** button will render the plain text `.Rmd` document into a variety of different output options.



The `.Rmd` files also give us the ability to document our intentions, write and execute code, and interpret and explain the results. After we've outlined (and revised) our thought process, we can go about organizing the code in more efficient ways to carry out our intentions.


Functional code chunks

We like to think of the `.Rmd` document like a stack of two sheets of paper, and each piece representing a different file type (`.R` and `.md`). When you lay the markdown file on top of the `.R` script, you get an Rmarkdown file.



R code chunk

This combination of files gives us a syntax for formatting our text (stuff like *italic*, **bold**, code, etc.), and a functional code script we can get access to by inserting code chunks. For example, in the `README.Rmd` file we can type directly onto the paper using the markdown syntax. But if we want to run some R code, imagine tearing a little hole in the markdown paper, and revealing an R script underneath.



R code chunk

These code chunks allow us to run R code in-between the markdown text.

R Code chunk options & labels

Code chunks come with a long list of options (feel free to experiment with all possible combinations found [here](#) and [here](#)). The most common are `echo`, `eval`, and `include`.

- `echo` = *show the code in the output?*
- `eval` = *run the code chunk?*
- `include` = *put the results from the code in the output?*

We will show a few examples of these with labels below:

```
```{r run-show-nothing, include=FALSE}  
runs and shows nothing
```
```

```
```{r run-show-code, eval=TRUE, echo=TRUE}  
runs code and shows results
```
```

```
```{r run-hide-code, eval=TRUE, echo=FALSE}  
runs code, shows results, but hides code
```
```

```
```{r no-run-show-code, eval=FALSE, echo=TRUE}  
shows code, but doesn't run code
```
```

The code chunks above are labeled, and we recommend *labeling all code chunks*. We've found it forces us to think through the analysis as a series of operations, and reduces random calculations spread throughout the document.

The setup code chunk

After the `YAML` header, our setup chunk tells us what we're going to be doing with the code, text, figures, and output in this `README.Rmd` file.

Our setup chunk does the following:

Project folders: We check to see if the project has image or data folders, and if not, it creates them.

```
# create image folder ----  
if (!file.exists("figs/")) {  
  dir.create("figs/")  
}  
# create data folder ----  
if (!file.exists("data/")) {  
  dir.create("data/")  
}
```


Chunk options: We've included `echo=TRUE`, which means the code will print in the output file. The `tidy=FALSE` makes sure the code doesn't get reformatted when the document gets rendered. We set both options to their default values, but we've included them below so that you can see more examples.

The `size = "small"` and `fig.path = "figs/"` are used to change the size of the printed code, and the location of any output visualizations.

```
knitr::opts_chunk$set(  
  echo = TRUE, # show all code  
  tidy = FALSE, # cleaner code printing  
  size = "small", # smaller code  
  fig.path = "figs/") # where the figures will end up
```

Knit options: The `knitr::opts_knit$set()` function gives us the ability to stipulate options for what happens when we render the document (knitted and rendered are somewhat synonymous).

```
knitr::opts_knit$set(  
  width = 78)
```

Base options: the final setting tells RStudio.Cloud how we want the tables to print (25 rows, and 78 columns).

```
base::options(tibble.print_max = 25,  
  tibble.width = 78)
```

Why 78? The standard code file is 80 columns across, which is the length of a punch card (read more [here](#)). We go two columns less than 80 (to give a little wiggle room).

Running code chunks

To run the code in your document, we have a few options. First, we can use the same keyboard shortcuts we use for executing code in the `.R` scripts (`ctrl+enter` or `cmd+return`).

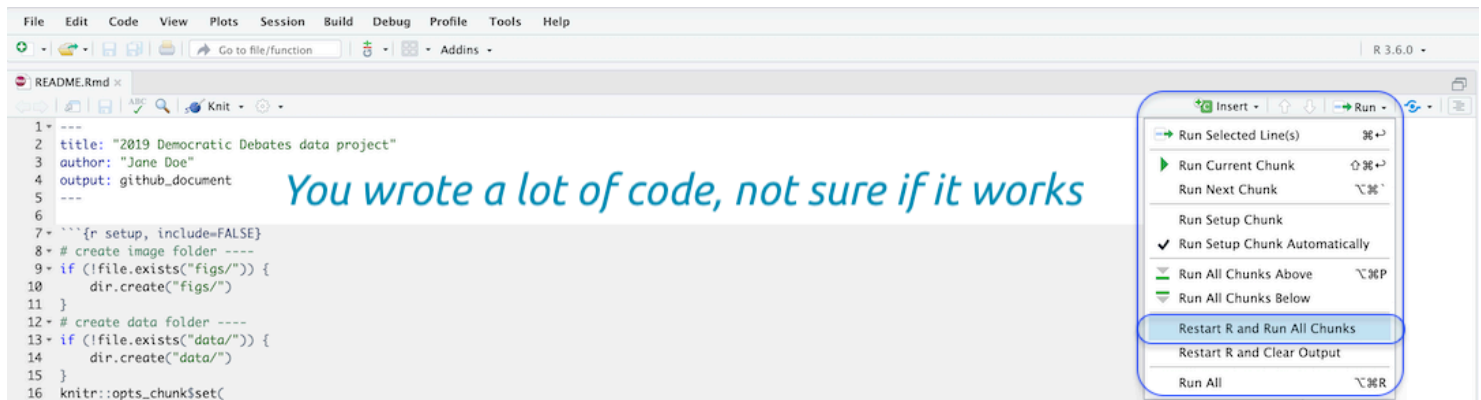
The second option is the small green play icon on the far right-hand side of the code chunk. Running code one chunk at a time is a great way to test what the output of a graph looks like, or to see the data table from a join.

```
```{r setup, include=FALSE}  
create image folder ----
if (!file.exists("figs/")) {
 dir.create("figs/")
}
```



**Click to run!**

The third option is we run a series of chunks (or the entire document) using the drop-downs at the top of the `.Rmd` file. We use this option if we've added a few chunks, and we want to make sure everything works together.



Don't worry if all of this is confusing! We will be creating our very own `README.Rmd` for this project, so you'll get some practice!

## Start with a README.Rmd

The `README.Rmd` file we've provided for this project is in the **Source** pane, and we can start by looking at the contents of the file (look at the YAML header, setup chunk, and packages).

As we've stated before, we won't be going through the code collects and downloads the data for this project. We'll provide links to resources where you can get all that information, but we're going to focus on the *workflow for using data to create tables, images, and products*.

Hit the enter/return key until we have a few lines of padding beneath the packages code chunk. Our cursor should be around line 37.

We will begin this document with the following text under our first line header (#).

### # Motivation

We read an interesting article on ``fivethirtyeight`` about the democratic debates in June of 2019. In particular, the image below displays how voters had changed their minds after watching the candidates.

As the text above states, we saw something cool, and we're going to see if we can recreate or verify some of the information. The source of this inspiration comes from [this fivethirtyeight article](#).

**Including an image in the .Rmd file:** If you want to include an image in your `README.Rmd` file, read about the `knitr::include_graphics()` function by entering `??include_graphics` in the **Console** pane.

## Data sources

Now that we have an idea why we're here, we want to start exploring the data. We're going to document all of our steps in the `README.Rmd` file.

We will start by importing all of the data for this project. We've put the raw data files in the `data/raw/` folder. **Always leave raw data in its own folder, and never alter the raw data files (read more [here](#)).**

In our `README.Rmd` file, we'll create a new level-two header (##) called 'Data files' and a code chunk labeled "locate-data" (it should look like the image below).

```

41 ▾ ## Data files
42
43 Below are the data files sorted by size:
44
45 ▾ ```{r locate-data}
46
47 ```
48

```

## File management with the fs package

To locate the files in `data/raw/` folder, we like to use the `fs` package. `fs` stands for ‘file system’, and this package gives us the ability to navigate our project folders and files.

Check out the package website [here](https://r-lib.github.io/infrastructure/rpackages/packages/fs/) for a full description and examples. `fs` is loaded as part of the `tidyverse`.

Start with a folder tree of the data folder.

```

where are the data?
fs::dir_tree("data")
data
└─ processed
└─ raw
└─ 538
└─ 2019-07-06-Cand538Fav.csv
└─ google-trends
└─ 2019-07-10-Dems2020Night1Group1.rds
└─ 2019-07-10-Dems2020Night1Group2.rds
└─ twitter
└─ 2019-07-06-Night01Tweets.rds
└─ 2019-07-06-Night01TweetsRaw.rds
└─ 2019-07-06-Night01TweetsUsers.rds
└─ 2019-07-06-Night02Tweets.rds
└─ 2019-07-06-Night02TweetsRaw.rds
└─ 2019-07-06-Night02TweetsUsers.rds
└─ wikipedia
└─ 2019-07-10-WikiDemAirTime01Raw.csv
└─ 2019-07-10-WikiDemAirTime02Raw.csv
└─ 2019-07-25-PollingCriterionRaw.csv

```

The raw data are in four separate folders, each representing the data source (538, google-trends, twitter, Wikipedia).

## Determine the size of the data files

Size can be a significant impediment to getting your work done quickly, so it's best to determine the size of the raw data files before importing.

The code chunk below tells us how big each file is and the folder in which its located.

```
how big are each of these files?
fs::dir_info(path = "data", recurse = TRUE) %>%
 # only files
 filter(type == "file") %>%
 group_by(path) %>%
 # sort by size
 tally(wt = size, sort = TRUE)

A tibble: 12 x 2
path n
<fs::path> <fs::bytes>
1 data/raw/twitter/2019-07-06-Night02Tweets.rds 7.99M
2 data/raw/twitter/2019-07-06-Night02TweetsRaw.rds 7.87M
3 data/raw/twitter/2019-07-06-Night01Tweets.rds 6.92M
4 data/raw/twitter/2019-07-06-Night01TweetsRaw.rds 6.83M
5 data/raw/twitter/2019-07-06-Night02TweetsUsers.rds 1.9M
6 data/raw/twitter/2019-07-06-Night01TweetsUsers.rds 1.65M
7 data/raw/google-trends/2019-07-10-Dems2020Night1Group2.rds 111.47K
8 data/raw/google-trends/2019-07-10-Dems2020Night1Group1.rds 109.52K
9 data/raw/wikipedia/2019-07-25-PollingCriterionRaw.csv 722
10 data/raw/538/2019-07-06-Cand538Fav.csv 581
11 data/raw/wikipedia/2019-07-10-WikiDemAirTime02Raw.csv 202
12 data/raw/wikipedia/2019-07-10-WikiDemAirTime01Raw.csv 191
```

The output tells us the twitter data files are the largest (7.99M). All of these files are small enough for RStudio.Cloud to handle, though.

## Loading the data into R

We can import the data using the 01-import.R file in the code folder. Feel free to open this file and examine its contents in the **Source** pane. We will use the `base::source()` function to run all of the code in the 01-import.R file.

```
import data using 01-import.R file
source("code/01-import.R")
```

The 01-import.R file loads all of the data files into the RStudio.Cloud session. We can verify this by examining the contents of the environment with `base::ls()`.

```
base::ls()
```

The output tells us the following objects are in our working environment.

```
[1] "google_data_files"
[2] "google_data_path"
[3] "GoogleData"
[4] "GSheetCand538Fav"
[5] "GTrendDems2020Night1G1"
[6] "GTrendDems2020Night1G2"
[7] "twitter_data_files"
[8] "twitter_data_path"
[9] "twitter_users_data_files"
```

```
[10] "TwitterData"
[11] "TwitterUsersData"
[12] "wiki_data_files"
[13] "wiki_data_path"
[14] "WikiData"
[15] "WikiDemAirTime01Raw"
[16] "WikiDemAirTime02Raw"
[17] "WikiPollCriterionRaw"
```

We can also see these objects in the **Environment** pane.

The screenshot shows the RStudio interface with the Environment pane on the right. The Environment pane is divided into two sections: 'Data' and 'Values'. The 'Data' section contains objects like GoogleData, GSheetCand538Fav, GTrendDems2020Night1G1, GTrendDems2020Night1G2, TwitterData, TwitterUsersData, WikiData, WikiDemAirTime01Raw, WikiDemAirTime02Raw, and WikiPollCriterionRaw. The 'Values' section contains objects like google\_data\_files, google\_data\_path, twitter\_data\_files, twitter\_data\_path, twitter\_users\_data\_files, wiki\_data\_files, and wiki\_data\_path. Annotations with arrows point to these sections, stating 'Data = tibbles & Lists' and 'Values = vectors'.

The *Global Environment* holds all of our imported data.

## Documenting changes to our project with Git

So, we've imported some data into the RStudio.Cloud session, but we need to make sure we're keeping track of the changes to our files.

We can do this by checking our `git status` in the **Terminal** pane.

```
$ git status
On branch master

Initial commit

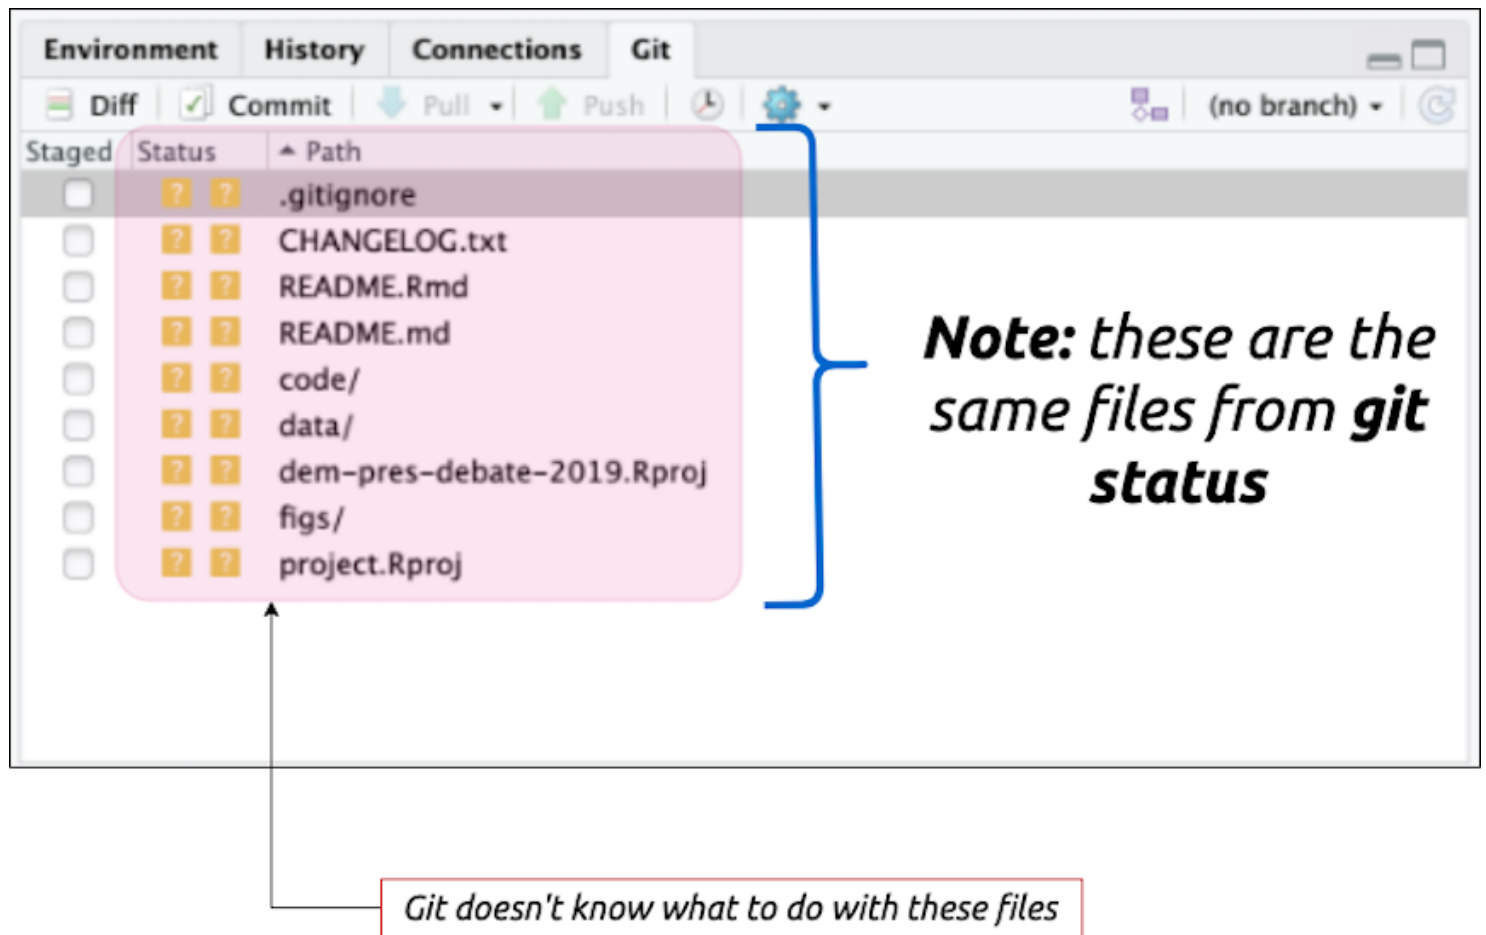
Untracked files:
 (use "git add <file>..." to include in what will be committed)

.gitignore
CHANGELOG.txt
README.Rmd
```

```
README.md
code/
data/
dem-pres-debate-2019.Rproj
figs/
project.Rproj
```

nothing added to commit but untracked files present (use "git add" to track)

Git is telling us we have quite a few new files in our project (which we expected), but that to git to pay attention to them, we need to add them. We can do this in the **Git** pane RStudio provides (it's by the **Environment** pane).



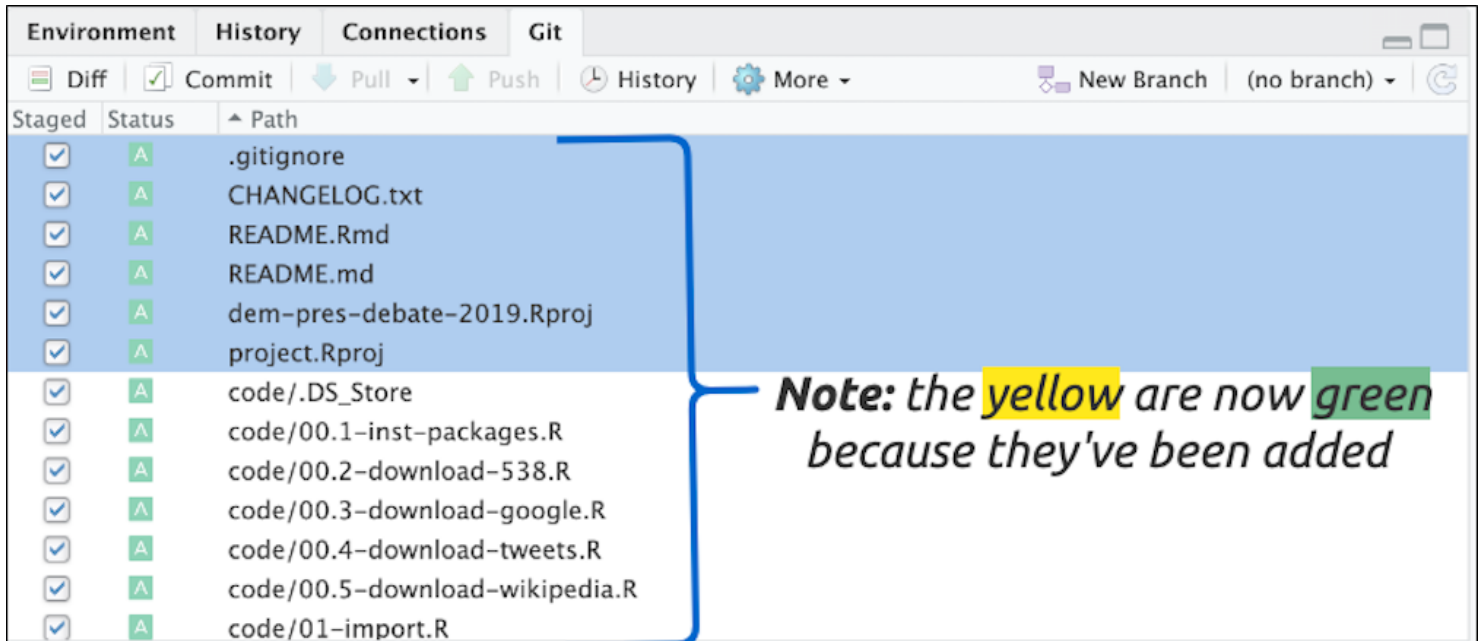
The screenshot shows the RStudio interface with the **Git** pane active. The pane has tabs for **Diff**, **Commit**, **Pull**, and **Push**. Below these tabs is a table with columns **Staged**, **Status**, and **Path**. The **Status** column shows orange question marks for untracked files. The **Path** column lists the files: `.gitignore`, `CHANGELOG.txt`, `README.Rmd`, `README.md`, `code/`, `data/`, `dem-pres-debate-2019.Rproj`, `figs/`, and `project.Rproj`. A blue bracket on the right side of the table points to the **Status** column with the text: **Note: these are the same files from git status**. A red box at the bottom of the pane contains the text: *Git doesn't know what to do with these files*, with an arrow pointing to the **Status** column.

Staged	Status	Path
<input type="checkbox"/>	??	.gitignore
<input type="checkbox"/>	??	CHANGELOG.txt
<input type="checkbox"/>	??	README.Rmd
<input type="checkbox"/>	??	README.md
<input type="checkbox"/>	??	code/
<input type="checkbox"/>	??	data/
<input type="checkbox"/>	??	dem-pres-debate-2019.Rproj
<input type="checkbox"/>	??	figs/
<input type="checkbox"/>	??	project.Rproj

**Note: these are the same files from git status**

*Git doesn't know what to do with these files*

We will click on each checkbox next to the files in the **Git** pane. Don't be alarmed if this list expands to include all the subfiles and folders.



Now we can re-check `git status` to see what just happened.

```
$ git status

On branch master

Initial commit

Changes to be committed:
 (use "git rm --cached <file>..." to unstage)

 new file: .gitignore
 new file: CHANGELOG.txt
 new file: README.Rmd
 new file: README.md
 new file: code/.DS_Store
 new file: code/00.1-inst-packages.R
 new file: code/00.2-download-538.R
 new file: code/00.3-download-google.R
 new file: code/00.4-download-tweets.R
 new file: code/00.5-download-wikipedia.R
 new file: code/01-import.R
 new file: code/02-wrangle.R
 new file: data/.DS_Store
 new file: data/processed/.DS_Store
 new file: data/raw/.DS_Store
 new file: data/raw/538/2019-07-06-Cand538Fav.csv
 # omitted
```

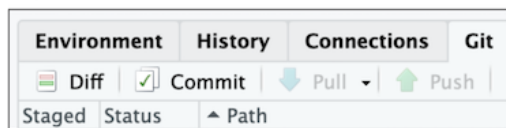
This long list of files shows all the contents have been added and are ready to be committed.

## Committing the changes

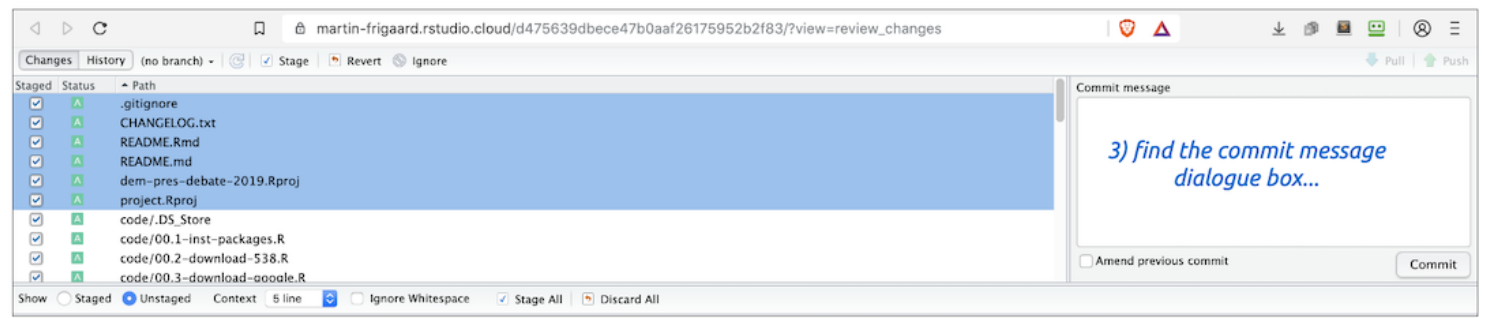


If we click on the *Commit* icon in the **Git** pane, this will open a new window in RStudio.Cloud. In this window, we will enter a commit message (“First commit!”), and click the *Commit* button. We’ve outlined this entire process in the schematic below:

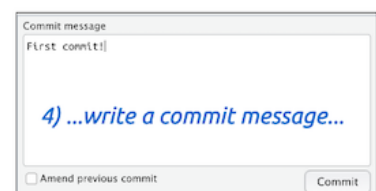
1) Click on the 'Commit' icon



2) This will show a new window, where we will see our added files




3) find the commit message dialogue box...



4) ...write a commit message...

5) click 'Commit'



6) Then click 'Close'

## A quick git terminology overview

Below are some commonly used terms/commands associated with Git and Github.

- init** - the command `git init` is used to initialize a new git repository (it tells Git to start tracking changes in this directory).
- status** - Git status will tell you what you’ve done and what is happening. You can check the status of a git repository with `git status` (use this liberally).
- add** - In order for Git to keep track of the changes we make to files, we have to tell it which files to pay attention to. We can do this using `git add`. The `git add -A` tells git to stage *ALL* the files that are in an initialized repo.
- commits** - commits are the staple in Git/Github the workflow. Commits are what Git uses to track the changes you’ve made to files or folders. Commits are confusing because they can be nouns (“I’m creating a commit with these changes”) or verbs (I am going to commit these changes to my project”).

To quote David Demaree,

*“Semantically, each commit represents a complete snapshot of the state of your project at a given moment in time; its unique identifier serves to distinguish that state from the way the files in your project looked at any other moment in time.”*

**repository** - repos are the files and folders in your project and all the changes you make while working on them. On your local computer, a repository can exist in a folder you initialize a repository in (see below). On Github, a repo has the following structure: `https://github.com/<username>/<repository_name>`.

**clone** - this command copies all files and changes into a new working directory from a remote, initializes (`init`) a new Git repository, and it adds a remote called `origin`.

**diff** - this is how Git shows differences between files. Read more about how changes are formatted/displayed [here](#).